

C-MCPN: The Clingo McCulloch-Pitts Network Editor

An Answer-Set Programming Framework for Artificial Neural
Networks

First Author¹, Second Author²
The Department Mathematical Sciences
University of Redacted
Location, ST, Zip

¹first email

²second email

Abstract

As Artificial intelligence and generative AI continue to engage the public in new and profound ways, there is growing interest in developing more efficient underlying technologies. In this paper we explore the use of efficient Answer Set Programming (ASP) languages to model and conduct inference over artificial neural network (ANN) architectures. As a supportive tool for our own reasearch we developed C-MCPN (Clingo McCulloch-Pitts Network) Framework and Editor, a graphical utility for representing and conducting inference over ANNs, with underlying representations in the ASP language, Clingo. In this paper we detail our underlying delcarative neural representations, and introduce C-MCPN as a new utility for general and educational use.

1 Introduction

Artificial Intelligence (AI) systems and their applications have been of significant consequence over recent decades. This has been especially evident in recent years abnormally significant awareness of AI in the general public. In particular, these include techniques from machine learning (ML) like the artificial neural networks (ANNs) that power contemporary generative AI (GEN-AI) systems like large language models (LLMs).

These AI systems have shown considerable promise, but the underlying technologies come with significant drawbacks. Concerns include the growing

footprint of data centers, which has garnered public interest, the worsening use of finite resources such as fresh water and electricity, and the potential for substantial environmental and health impacts [Han2025, Siddik2021, Li2023].

Conversely, Answer Set Programming (ASP) languages live at a different end of the research spectrum in terms of their efficiency. They are designed to solve NP-Complete and NP-Hard problems in a fast and efficient manner on limited hardware. Our current work involves exploring the use of ASP languages to accelerate the processes of learning and inference. The eventual goal of our research is to reduce the resources needed for these technologies.

In this paper we introduce our ANN representations in the Clingo language. We also introduce the C-MCPN: The *Clingo McCulloch-Pitts Network Editor*. The C-MCPN is a tool that we developed to assist in the exploration of ASP-based networks. It allows a user to model the network architecture and conduct inference, with underlying representations and inference guided by declarative software written in Clingo. Specifically we restrict our attention to McCulloch-Pitts Networks (MPNs) as the target ANN type [McCulloch1943].

2 Background

In this section we will expand upon the preliminaries of our work, starting with a short discussion regarding the potential environmental impacts of current generation GEN-AI systems. We then provide a brief background on both ANNs and the ASP dialect Clingo, motivating our work on C-MCPN.

2.0.1 Environmental Impacts

The introduction of big data, chiefly by scraping the broader Internet, changed the landscape for AI. It allowed for statistical-based models to demonstrate both power and storage capacity, which require massive amounts of data for training purposes [Goodfellow2016]. This incredible need for information highlights an important drawback in the form of *data efficiency* [Cropper2020]. The lack of efficiency then translates to a significant amount of necessary resources for both initial training and general use.

As an example, consider the cost to train a single model—ChatGPT-3. GPT-3 was reported to require around 1.3 million kWh *just* to be trained [DeVries2023]. With some quick, back-of-the-envelope calculations, we can put this into perspective. In 2022 the average American household used approximately 10,791 kWh [EIA2024], and by dividing these out we find that

the model’s energy use was proportional to:

$$1,300,000 \text{ kWh} \times \frac{1 \text{ household}}{10,791 \text{ kWh/year}} \approx 120.5 \text{ households.}$$

This is a remarkably over-simplified example, ignoring many other contributing factors. E.g., We ignore any *indirect* costs to train the model and produce the energy, the requirements to develop the infrastructure, and any resources required to *deploy and run* the model. However this does highlight both the significant resources needed to employ these technologies, the dire need for more efficient, sustainable AI and ML technologies [DeVries2023]. Unfortunately reducing the environmental impacts of such technology is still an open problem.

2.1 Artificial Neural Networks (ANNs)

ANNs were originally inspired by the structure and function of biological neurons. Their basic building blocks, the artificial neurons, can be traced back to the work of McCulloch and Pitts in the 1940s [McCulloch1943]. These early models mimicked the behavior of biological neurons by using a simple thresholding function that determines whether a neuron will fire or not. The name for the C-MCPN was inspired by this early work. Despite the name and our examples, the neural models are sufficiently extensible to represent neurons beyond MCP neurons.

The neurons can be connected together to form a network, whose structures can vary widely from simple knowledge graphs to feedforward networks to more complex recurrent architectures [Zhou2020]. For our work, the feedforward neural network was chosen for its simplicity and wide use in the field. It consists of layers of interconnected neurons that process data in a single forward direction [Goodfellow2016], shown in Figure 1. These models can learn through a training process where weights are adjusted based on error, but notably their underlying structure is similar to that of the original McCulloch-Pitts model.




Figure 1: A diagram of a simple feedforward neural network

Computation is accomplished through a similar network structure. Initial values feed into the network and propagate according to the underlying network architecture, and the functions encoded for individual neurons to fire. Solutions are derived from values at the opposite end of the network.¹

2.2 Answer Set Programming (ASP)

Answer Set Programming serves as an amalgamation of Prolog-style quantified logic with powerful advances in satisfiability (SAT) solvers that occurred throughout the 1990s and into the 2000s. In this section we will discuss the basic syntax and semantics of the ASP language, Clingo, which serves as a technological foundation of C-MPCN [Gebser2019].

2.2.1 Syntax and Semantics

The atomic units in Clingo are *literals*, which describe ideas and *facts* that serve as statements that are held to be true. The basic syntax is the application of a *property* to a literal, `property(atom)`. E.g.,

```
red(apple).
```

is a complete statement (and program), attributing a property of red to an idea named `apple`.

Further knowledge is *inferred* using aggregate statements known as *rules*, taking the form: `consequent :- antecedent`. This mirrors an *implication* from formal logic, as the `:-` symbol is read equivalently to $\leftarrow$. In such a statement the consequent *must* be true whenever the antecedent is true. E.g.,

```
eat(Fruit) :- red(Fruit), ripe(Fruit).
```

describes a preference over what fruit will be eaten. I.e., “If a fruit is red and ripe, then we eat the fruit.”

In this case rather than the literal `apple`, we use a *variable* to allow for inference. The `Fruit` can be *any* contextually-relevant literal in the program, so long as it possesses the required properties.

3 Methodology

Combining the efficiency of ASP languages with the representational power of ANNs is a non-trivial task. Not only does it require a novel representation

¹In reality, such networks are *learned* from data due to the efforts required to construct them by hand. While efficient structure and parameter learning are an inevitable goal of this line of research they are outside of the scope of this paper.

of the ANN structure, but Clingo has limited floating-point support. This requires clever simplifications of the activation functions and gradient descent to use them over integers. Additionally, large ANNs can cause combinatorial explosions in the search space, making grounding the programs difficult. The C-MCPN Framework and Editor were developed to address some of these challenges, providing a structured way to represent and work with logical networks in Clingo.

3.1 Domain Selection

To build and test the framework the illustrative domain of digital logic for hardware development was selected. Its simplicity and familiarity make it an ideal domain for experimentation.

For our example network, a 1-bit adder seen in Figure 2, we used neurons with three types of well established logic gates: AND, OR, and XOR. In Table 1 we show the neuron definitions of these gates, and their truth tables. Using a combination of these and more gates, one is able to construct complex logical operations².

AND			OR			XOR		
$g(\vec{x}) = x_1 \cdot x_2$			$g(\vec{x}) = x_1 + x_2$			$g(\vec{x}) = x_1 - x_2 $		
x_1	x_2	$f(\vec{x})$	x_1	x_2	$f(\vec{x})$	x_1	x_2	$f(\vec{x})$
T	T	T	T	T	T	T	T	F
T	F	F	T	F	T	T	F	T
F	T	F	F	T	T	F	T	T
F	F	F	F	F	F	F	F	F

Table 1: Neuron Definitions and Truth Tables.

Other gates including negation will use similar constructions that follow from the above.

3.2 Design Approach

The representation of an artificial neuron requires both *structure* designation, and activation function annotation. When activated the neuron will propagate a signal to its output given the provided input signals. Its value then further propagates through the network along edges and activating other neurons.

An AND gate is used to illustrate how we represent an artificial neuron:

²An elegant and simple proof of Clingo’s Turing completeness falls out of the C-MCPN Framework.




Figure 2: A screenshot of a 1-bit adder network in the C-MCPN Editor.

```
% gate(TYPE, OUTPUT, INPUTS)
gate("AND", Out, In1, In2) :- values(In1), values(In2),
                                Out = (In1 * In2).
```

The *structure* of the neuron is indicated by the parameters: in the case of an AND gate this requires two input parameters and a single output. The *activation function* is encoded directly in the rule structure using the arithmetic equivalent provided in Table 1.

The neuron *connectivity* specified in the gate definitions can then be used to forge connections of a network. A connection from the 1-bit adder sample network is built as:

```
% neuron(type, unique_id)
neuron("INPUT", "INPUT_1").
neuron("XOR", "XOR_1").

% edge(source_id, target_id)
edge("INPUT_1", "XOR_1").

% val(neuron_id, value).
val("INPUT_1", 1).
```

Finally, the value inference program is used to store the possible values, activate the neurons, and propagate the values through the network. The program for a 2-input gate can be built as:

```
%% Value domain
values(0;1).

% Edges carry values from input to output
edge(In, Neuron_id, Val) :-
    edge(In, Neuron_id),
    neuron(_, Neuron_id),
```

```

    neuron(_, In),
    val(In, Val).

% 2 input gates
val(Neuron_id, Val) :-
    gate(Type, Val, Val1, Val2),
    neuron(Type, Neuron_id),
    edge(In1, Neuron_id, Val1),
    edge(In2, Neuron_id, Val2),
    In1 != In2.

%% Each neuron has one value
{ val(Neuron_id, Val) : values(Val) } 1
    :- neuron(_, Neuron_id).

%% Show values
#show val/2.

```

The rule carrying the values along the edges is vital to the functioning of the network. Without it, the network gets "stuck" and is unable to infer what the values should be—a key insight discovered during development.

4 C-MCPN

To better explore the ANNs implemented within Clingo we developed the C-MCPN software package to automate our workflow. The system itself is designed around a standard Unix terminal workflow and is extended with a graphical user interface (GUI).

The core functionalities of C-MCPN are built on the ASP language Clingo, as described in the preceding sections. The editor and interface were constructed in Python using the DearPyGui library, allowing us to take advantage of its node-based editing capabilities.

For node reorganization in the editor we used a modified version of the Sugiyama algorithm for layered graph drawing, taking just the layer assignment step and a simplified node positioning step [Sugiyama1981].

4.1 C-MCPN Framework

From what the authors found in the literature, the C-MCPN Framework is a novel approach to neural network representation within a logical programming paradigm like ASP. The networks themselves have relatively simple representations that can be leveraged to create complex operations. The framework, which can be seen in Figure 3, uses three main program types, all of which

are contained in the C-MCPN/ directory of the project. The content in the gate definitions program defines the behavior of each gate, but every default gate we included also has a special header used to send meta-information to the C-MCPN Editor.




Figure 3: A visualization of the programs involved in the C-MCPN Framework.

4.2 C-MCPN Editor

The `app/` directory of the project contains the source code for the C-MCPN Editor, which acts as a lightweight interactive development environment (IDE) for logical networks. Its main features include the ability to add/remove neurons, connect/disconnect them, copy/paste them, and run inference on the network. Additionally, the editor provides a live preview of the ASP code generated from the network structure, allowing users to see how their changes affect the underlying logic program.

User created networks are saveable to JSON and ASP format, and loadable from JSON format. This facilitates creation of complex ANNs, and saving the network to ASP format allows direct use of the C-MCPN Framework from the terminal instead of the IDE.

A screenshot of the editor can be seen in Figure 4, showcasing the sample 1-bit adder network.

4.3 Development

Claude (Anthropic) was used as a programming assistant throughout development of the editor, assisting with module integration, debugging, and some code generation. All design decisions, research direction, and domain-specific logic (ASP, gate definitions, network structure) were conscious decisions by the authors. Each line of code generated by Claude was reviewed, edited, and approved by the authors before being added to the codebase. In order to maintain a supervisory-level of control the agent was not given direct access to the codebase or the repository.




Figure 4: Screenshot of the C-MCPN Editor, showing a sample 1-bit adder network.

Use of Claude was employed only in the above described sections in order to accelerate the development process under circumstances of limited resources.³

5 Conclusions & Future Directions

The primary results of our work are a representation for ANNs in a declarative, ASP language, as well as the C-MCPN Framework and Editor. Results on efficiency and scalability, our primary motivation, are left for future work, as they matter more for the adaptive, learning-capable network framework we are currently developing. Additionally, we hope to continue adding features into the future regarding inference over more network structures, including recurrent networks.

Both tools are open source and available at **cite anonymized repository here**.

³The authorship of this paper is without the assistance of any AI system, including planning, writing, editing, and repository maintenance.